

Object Drive 1.0

System Admin Guide

(SAG)

Refer to System Installation Procedures Guide for Configuration

Last Updated: September 13, 2019

Master Encryption Key	3
Backing up Encryption Key	3
Changing the Encryption Key	3
FIPS 140-2	7
Verifying the Artifacts Used to Compile	7
Compare with MD5SUM from Google Bucket	8
Compare with SHA256SUM	8
Verifying Binary Built with Go is using BoringCrypto	8
Object Drive Reporting	9
Command Line Flag	9
Logging	10
API Documentation	10
RPM Filename	10
Docker Images	11
Useful SQL Queries	12
See when schema was created	12
Check for presence of a function or procedure by name	12
List all tables in database, and space used	12
Size of the database	12
Count of Objects by Document Size Ranges	13
Count of Objects by Period of Time	13
Fixing User Authorization Object Cache from Snippets	14
Removing Duplicate User Cache Records	14
Brute Force Clearing User Cache and ACM Associations	15
Identifying and Fixing Faulty ACMS	16

File Cache Management	19
Cache Purge Cycle	19
Cache Purge Iteration	20
File Purge Visit	21

Changes

Author	Date	Description
Lucas Moten	2019-09-13	Added sql query to get Count of Objects by Period of Time

Master Encryption Key

The master encryption key is configured during initial setup as described in the System Installation Procedures Guide (SIPG)

For convenience, here is the description of this environment variable

OD_ENCRYPT_MASTERKEY
since v1.0

The secret master key used as part of the encryption key for all files stored in the system. If this value is changed, all file keys must be adjusted at the same time. If you don't set this, the service will shut down. Note that if a token.jar is installed onto the system, we can use the encrypted indicator format like `ENC{...}`

Every instance of Object Drive within the cluster **MUST** be using the same key.

Backing up Encryption Key

The encryption key should be backed up to a place where it can be restored as necessary in the event that all instances of Object Drive server are terminated and deleted. This may be maintained by an ISSO or other individual in charge of managing secrets in escrow.

Changing the Encryption Key

Periodically, the master encryption key may need to be changed.

1. **Backup the database!** This action may be long running and is irreversible. **You should always make a snapshot or backup of the database before applying a system wide change like this or else you run the risk of losing all of the data.** Typical disaster recovery plans (DRP) also call for verifying the backup, but that is outside the scope of this procedure.
2. **Stop all running instances of Object Drive.** To prevent new objects being created, or changes from the service transpiring during the update, **all services should be stopped.**

3. **Using the credentials to access the database**, login to the database using the mysql client.
4. **Set database to read only mode.** This step should only be performed on instances running version 1.0.2 (June 30, 2017) or newer. The services will not attempt to make changes if the value of schemaversion stored in the dbstate table is different than the internally expected version as of v1.0.2.

```
update dbstate set schemaversion =  
concat(replace(schemaversion, 'READONLY', ''), 'READONLY');
```

5. **Verify that the rotate_keys procedure exists.** The following SQL statement can be typed in to validate existence of the procedure. A result of 1 indicates the procedure exists.

```
select count(0) from information_schema.routines  
where routine_schema = database() and routine_name = 'rotate_keys';
```

If it does not exist, then you may create it as follows

```
CREATE PROCEDURE rotate_keys(IN oldMasterPass VARCHAR(255), IN newMasterPass  
VARCHAR(255), IN entropy VARCHAR(255))  
BEGIN  
  declare exit handler for sqlexception  
  begin  
    rollback;  
  end;  
  declare exit handler for sqlwarning  
  begin  
    rollback;  
  end;  
  start transaction;  
  update object_permission p set encryptKey = unhex(  
    bitwise256_xor(  
      new_keydecrypt( newMasterPass, hex(p.permissionIV)),  
      bitwise256_xor(  
        new_keydecrypt(oldMasterPass, hex(p.permissionIV)),  
        hex(p.encryptKey)  
      )  
    )  
  );  
  update object_permission p set permissionMAC = unhex(new_keymac(newMasterPass,
```

```
p.grantee, p.allowCreate, p.allowRead, p.allowUpdate, p.allowDelete, p.allowShare,  
hex(encryptKey));  
commit;  
END;
```

4. **Make note of the current encryption key.** For a typical installation, this value can be found in the file `/opt/services/object-drive-1.0/env.sh` named `OD_ENCRYPT_MASTERKEY`.
5. **Determine a new value to use as the master encryption key.** As a convenience, you can generate a hash from the command line, or a simple hex of the current time in MySQL as follows. Whatever approach you take to determine your new encryption key, make sure it follows guidelines of strong and secure passwords outlined by your organization. **The ISSO should backup the new encryption key.**

```
select hex(current_timestamp());
```

A sample run of the above produces output similar to the following:

```
323031382D30382D30372031383A33383A3133
```

If you want something longer, you can use the following which includes some additional database specific functions and trims the left and right to produce a desired length.

```
select lcase(reverse(right(left(concat(hex(current_timestamp()), hex(database()),  
hex(session_user()), 80), 60)));
```

A sample run of the above produces output similar to the following:

```
e223731304275637572646264616471646164756d65343a36333a3831302
```

6. **Execute the change.** You can call the procedure as follows replacing the `currentKey` and `newKey` with the string values identified in steps 4 and 5 respectively:

```
CALL rotate_keys('currentKeyFromStep4', 'newKeyFromStep5', 'reserved-for-entropy');
```

This operation may take some time to complete as it will update all permissions on all objects stored in this system. The larger the instance, the more time it will take to complete.

At present, there is no batch processing for this change, and all changes take place in a single transaction.

There is no indication of success or failure until attempting to access a file. If you provide the wrong key in step 4, the data may be irrevocably lost. It is critical to ensure that you've taken a backup in step 1.

7. **Start service instances.** If you stopped the object drive server instances, you may restart them after updating the OD_ENCRYPT_MASTERKEY value in `/opt/services/object-drive-1.0/env.sh` with the new master key derived in step 5. If you chose to put instances in read only mode, then you should shut down all instances now before making the change
8. **Perform a Test.** Now is the time to do a test to verify that you and/or other users can continue to access the streams of files you were able to before. If you are able to download a file and open it, then the change was successful.

However, if a file you were able to access previously now downloads as complete gibberish, then something went wrong with the key rotation. If this happens, it's recommended to restore the OD_ENCRYPT_MASTERKEY to the value it was before, and restore the database snapshot and re-attempt. If you continue to have problems, reach out to developers.

9. **Enable Write Access.** Now you can set the write access to the database by changing the schema version back to what it was before. This step is only necessary if you set object drive to read only in step 4.

```
update dbstate set schemaversion = replace(schemaversion,'READONLY','');
```

10. **Done.**

FIPS 140-2

From wikipedia

The Federal Information Processing Standard (FIPS) Publication 140-2, (FIPS PUB 140-2), is a U.S. government computer security standard used to approve cryptographic modules. The title is Security Requirements for Cryptographic Modules. Initial publication was on May 25, 2001 and was last updated December 3, 2002

Object Drive is built using the Go programming language, freely available from Google at <https://golang.org/dl/> or with BoringCrypto enabled versions, available here: <https://go-boringcrypto.storage.googleapis.com/>

Natively, the Go programming language does not support FIPS 140-2 compliance. Per the [National Institute of Standards and Technology \(NIST\)](#), compliance is through use of those modules that have been successfully vetted through the [Cryptographic Module Validation Program](#) (CMVP).

The Go programming language has branches, starting with Go version 1.8, that include alterations in the crypto and other packages to leverage the core of the BoringSSL codebase as a module named BoringCrypto. This module was validated through NIST and has [Certificate #2964](#) for software version [24e5886c0edfc409c8083d10f9f1120111efd6f5](#). The branch of go is '[dev.boringcrypto](#)'.

Verifying the Artifacts Used to Compile

The boringssl code for that software version may be downloaded from <https://commondatastorage.googleapis.com/chromium-boringssl-docs/fips/boringssl-24e5886c0edfc409c8083d10f9f1120111efd6f5.tar.xz> with a sha256 checksum of 15a65d676eeae27618e231183a1ce9804fc9c91bcc3abf5f6ca35216c02bf4da. This download and check is done within the Go programming language branches incorporating BoringCrypto as seen here: <https://github.com/golang/go/blob/dev.boringcrypto.go1.11/src/crypto/internal/boring/build/build.sh>. This compiles the module written in C, and [CGO](#) capabilities allows it to be called from Go. When a program is compiled using a version of Go that has BoringCrypto baked in, it makes callouts to the approved module in place of the native code to do the same.

Compare with MD5SUM from Google Bucket

The [gsutil utility](#), available as part of the google-cloud-sdk available from google can be used to report on the MD5sum of a file as stored in the google bucket, as well as calculate the same locally. For example, running this command:

```
gsutil stat gs://go-boringcrypto/go1.11b4.linux-amd.tar.gz
```

reports the md5 hash to be /yvx0VhhFQjMaeoc1ce/kg==, which is the base64 of the 128bit hash. The hex equivalent is ff2bf1d158611508cc69ea1cd5c7bf92 . This is the value you'd get if you downloaded the file using wget, and then ran md5sum on the file. You can also convert from base64 to hex like this

```
echo "/yvx0VhhFQjMaeoc1ce/kg==" | base64 -d | xxd -p
```

which will return the hex equivalent that md5sum will report.

```
ff2bf1d158611508cc69ea1cd5c7bf92
```

Compare with SHA256SUM

A sha256sum comparison can be done on the file as well, and the correct expected hash values are maintained on the RELEASES page for the dev.boringcrypto branch found here:

<https://go.googlesource.com/go/+dev.boringcrypto/misc/boring/RELEASES>

This is an HTML based resource from the Google git repository. For text, add ?format=text to the end, which will return it base64 encoded. To Get just the sha256sum part of a line from a particular release, the following one liner can be used:

```
curl -s https://go.googlesource.com/go/+dev.boringcrypto/misc/boring/RELEASES?format=text |  
base64 -d | grep go1.11b4.linux-amd64.tar.gz | awk '{print $5}'
```

For the file downloaded locally, running

```
sha256sum go1.11b4.linux-amd.tar.gz
```

It should produce the following as the result

```
d53417b2071af0104fbc15a957000bccdcb5bbc094df0401f67d51968f7f2e4e
```

Verifying Binary Built with Go is using BoringCrypto

Since compiling a go program creates a single binary to run on the architecture it is compiled for, it's useful to have ways to verify what actually makes up that binary.

If the Go programming language is available on the system, you can run the following command on a previously produced binary to see what symbols it has in use

```
go tool nm <binary-name>
```

For example, `go tool nm odrive`. This will produce a lengthy output of the symbols used. To narrow down to boring crypto, you can `grep` the output as follows

```
go tool nm odrive | grep "crypto/internal/boring._cgo"
```

This will reflect those functions that are running through the boring crypto module. Some expected output should appear similar to the following (truncated for brevity)

```
17f58f0 D crypto/internal/boring._cgo_925316647872_Cfunc_EVP_AEAD_CTX_open_wrapper
17f58f8 D crypto/internal/boring._cgo_925316647872_Cfunc_EVP_AEAD_CTX_seal_wrapper
17f5900 D crypto/internal/boring._cgo_925316647872_Cfunc__Cmalloc
17f5908 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_cbc_encrypt
17f5910 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_ctr128_encrypt
17f5918 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_decrypt
17f5920 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_encrypt
17f5928 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_set_decrypt_key
17f5930 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_AES_set_encrypt_key
17f5938 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_BN_bin2bn
17f5940 D crypto/internal/boring._cgo_925316647872_Cfunc__goboringcrypto_BN_free
```

```
...//snipped//...
```

Object Drive Reporting

Starting with v1.0.18, Object Drive will now report in a few different ways whether it was built with a version of Go that has the BoringCrypto module in place for crypto related calls

Command Line Flag

A command line flag "isfips" is added. It can be called as follows:

```
odrive1 isfips
```

This will indicate a FIPS 140-2 compliance check for the BoringCrypto module availability and exit.

¹ The name of the binary is by default 'odrive' within the development environment and docker images. For RPM deployed instances, the name of this binary is object-drive-1.0

FIPS 140-2 compliance check for BoringCrypto module

```
Built with Go runtime.version: go1.11b4
    base_go_version = go1.11
    boringcrypto_enabled = b
    boringcrypto_update_version = 4
```

The version of Go used to compile this binary uses BoringCrypto

If Go is available, you may also check whether the binary is using symbols from the module
`go tool nm odrive | grep "crypto/internal/boring._cgo"`

Sample Output from calling: `odrive isfips`

Logging

When started, logging output will also check the runtime.Version and report if its built using a version of the Go programming language that uses BoringCrypto. Here is a sample log entry

```
2018-11-02T11:28:26.608740536-04:00 INFO boring-crypto {"update": "4", "runtime.Version": "go1.11.5b4"}
```

API Documentation

The API documentation includes the version, build number, and build date at the top of each page. When the service is built with a version of Go programming language that uses BoringCrypto, it will also include the boring crypto update suffix to the version number.

- With BoringCrypto
 - **1.0.23b4** - The 'b4' indicates it was built with BoringCrypto update 4
- Without BoringCrypto
 - **1.0.23** - This indicates a version of Object Drive built without BoringCrypto

RPM Filename

Like the API documentation, if the service is built with a version of Go programming language using BoringCrypto, it will also include the boring update suffix in the RPM filename as part of the version. The following are examples with and without BoringCrypto respectively

- With BoringCrypto
 - `object-drive-1.0-1.0.23b4-3075.20190913.x86_64.rpm`
- Without BoringCrypto
 - `object-drive-1.0-1.0.23-3075.20190913.x86_64.rpm`

Docker Images

Docker images tagged for a version will have a suffix of "b" and a number indicating they include binaries built with BoringCrypto.

- With BoringCrypto:
 - deciphernow/**odrive-bc**:latest
 - docker-dime.di2e.net/dime/object-drive-server:**1.0.23b4**
 - docker-dime.di2e.net/dime/object-drive-server:**1.0.23b4-release-candidate**
- Without BoringCrypto
 - deciphernow/**odrive**:latest
 - docker-dime.di2e.net/dime/object-drive-server:**1.0.23**
- No guarantee of whether BoringCrypto is included or not:
 - docker-dime.di2e.net/dime/object-drive-server:**latest**

Use of a docker image without specifying the tag will result in the latest tag being used. This should never be done for a production environment, and should be considered unstable with no guarantee of what version it actually represents. Only tags for specific versions are supported.

Useful SQL Queries

See when schema was created

```
select create_time from information_schema.tables
where table_schema = database() and table_name = 'dbstate';
```

Check for presence of a function or procedure by name

```
select created, last_altered, routine_name
from information_schema.routines
where routine_schema = database() and routine_name = 'aacflatten';
```

List all tables in database, and space used

```
SELECT table_name, table_rows, data_length, index_length,
round(((data_length + index_length) / 1024 / 1024),2) "Size in MB"
FROM information_schema.TABLES where table_schema = database();
```

Size of the database

```
SELECT table_schema "Data Base Name",
sum( data_length + index_length) / 1024 / 1024 "Data Base Size in MB"
FROM information_schema.TABLES GROUP BY table_schema ;
```

Count of Objects by Document Size Ranges

```
select '0 - 5MB', count(0) from object
where contentsize >= 0 and contentsize < (5*1024*1024)
union select '5MB - 10MB', count(0) from object
where contentsize >= (5*1024*1024) and contentsize < (10*1024*1024)
union select '10MB - 25MB', count(0) from object
where contentsize >= (10*1024*1024) and contentsize < (25*1024*1024)
union select '25MB - 100MB', count(0) from object
where contentsize >= (25*1024*1024) and contentsize < (100*1024*1024)
union select 'Larger than 100MB', count(0) from object
where contentsize >= (100*1024*1024);
```

Here's an alternative to the above which also breaks down by object type and content type

```
select '0 - 5MB' streamsize, ot.name typename, o.contentType, count(0) qty from object o
inner join object_type ot on o.typeid = ot.id
where o.contentsize >= 0 and contentsize < (5*1024*1024) group by 1,2,3
union select '5MB - 10MB', ot.name, o.contentType, count(0) from object o
inner join object_type ot on o.typeid = ot.id
where o.contentsize >= (5*1024*1024) and contentsize < (10*1024*1024) group by 1,2,3
union select '10MB - 25MB', ot.name, o.contentType, count(0) from object o
inner join object_type ot on o.typeid = ot.id
where o.contentsize >= (10*1024*1024) and contentsize < (25*1024*1024) group by 1,2,3
union select '25MB - 100MB', ot.name, o.contentType, count(0) from object o
inner join object_type ot on o.typeid = ot.id
where o.contentsize >= (25*1024*1024) and contentsize < (100*1024*1024) group by 1,2,3
union select 'Larger than 100MB', ot.name, o.contentType, count(0) from object o
inner join object_type ot on o.typeid = ot.id
where o.contentsize >= (100*1024*1024) group by 1,2,3;
```

Count of Objects by Period of Time

You can select and group by the createddate or modified date of objects with varying granularity to determine the number of objects created during different spans of time. This is recommended to be analyzed on a monthly basis to help project data growth.

```
SELECT left(createddate,10), count(0) FROM object GROUP BY left(createddate,10);
```

Additional filters can be easily added to the where clause if you need to restrict the query by individual user, group, or parent folder.

Fixing User Authorization Object Cache from Snippets

If users are not able to see folders and files shared to them due to an error that could have occurred during caching when running a version of Object Drive prior to 1.0.13

```
update useraocache set sha256hash = '0', iscaching = 0 where iscaching > 0 and  
timestampdiff(minute, cachedate, now()) > 5;
```

You can also target a specific user as follows

```
update useraocache set sha256hash = '0', iscaching = 0 where userid in (  
  select id from user where distinguishedname =  
  'cn=test tester07,ou=people,ou=dae,ou=chimera,o=u.s. government,c=us'  
);
```

Note that the issue that can cause this may happen again until an upgrade to the service is performed that addresses deadlocks at the database tier. After either command type is executed, when a user accesses ODrive again, it will start fresh with attempting to establish the cache of ACMs they can access based upon their user profile

Removing Duplicate User Cache Records

Under certain scenarios, it's possible that a parallel request won't be properly detected by the service and this can result in multiple user cache records being added. A subsequent restart of the service will then leave affected users unable to access the system as a key handler expects only one record to be present and will roll up a failure as a result.

These duplicate records may be periodically cleared by executing the following SQL statement

```
delete e.*  
from useraocache e  
where id in (  
  select id  
  from (  
    select id  
    from useraocache  
    where userid in (  
      select a.userid  
      from (  
        select userid, count(0)  
        from useraocache  
        group by userid  
        having count(0) > 1  
      ) as a  
    ) and id not in (  
    select id  
    from useraocache  
    where userid in (  
      select a.userid  
      from (  
        select userid, count(0)  
        from useraocache  
        group by userid  
        having count(0) > 1  
      ) as a  
    )
```

```

select maxid
from (
  select ua.id, max(ua.id) maxid, hex(ua.userid)
  from useraocache ua inner join (
    select userid, count(0)
    from useraocache
    group by userid
    having count(0) > 1
  ) as a on ua.userid=a.userid
  group by ua.userid
) as b
)
) as c
);

```

Brute Force Clearing User Cache and ACM Associations

To fully clear the cache AO for a user, and any existing associations to ACMs, the following queries can be performed

```

delete from useraocachepart
where userid = any(
  select id from user
  where distinguishedname = 'cn=test tester10,ou=people,ou=dae,ou=chimera,o=u.s. government,c=us'
);

delete from useraocache
where userid = any(
  select id from user
  where distinguishedname = 'cn=test tester10,ou=people,ou=dae,ou=chimera,o=u.s. government,c=us'
);

delete from useracm
where userid = any(
  select id from user
  where distinguishedname = 'cn=test tester10,ou=people,ou=dae,ou=chimera,o=u.s. government,c=us'
);

```

To apply this for all users, simply reduce the where clauses

```

delete from useraocachepart;

delete from useraocache;

delete from useracm;

```

It will take a few seconds to rebuild the associations for a user when they first access the system. Be forewarned that doing this may result in unforeseen issues with concurrency as multiple users login/access Object Drive services in a small span of time.

Identifying and Fixing Faulty ACMS

An ACM is stored in multiple ways internal to Object Drive to improve search performance. A flaw in the way the data for ACM association was uncovered and known to affect 1.0.17, and 1.0.18, and likely several versions prior to that.

While 1.0.19b4 fixes the cause of this issue, it is prudent to check to see if there are any ACMS that may have been created in the system incorrectly.

To check for faulty ACMS run the following query

```
select id, flattenedacm from acm2 where id not in (select acmid from acmpart2);
```

If there are any rows returned, then the following script can be run to repair them. This will be included in a database migration as part of a future patch release.

```
DROP PROCEDURE IF EXISTS sp_Patch_20190225_fix_acmpart2_associations;
DELIMITER //
CREATE PROCEDURE sp_Patch_20190225_fix_acmpart2_associations()
BEGIN
    -- declare variables
    DECLARE vACMID int default 0;
    DECLARE vACMName text default '';
    DECLARE vSHA256Hash char(64) default '';
    DECLARE vKeyValuePart varchar(2000) default '';
    DECLARE vKeyID int default 0;
    DECLARE vKeyName varchar(255) default '';
    DECLARE vValueID int default 0;
    DECLARE vValueName varchar(255) default '';
    DECLARE vPartID int default 0;
    DECLARE c_acm2_finished int default 0;
    DECLARE c_acm2 cursor for SELECT id, flattenedacm
        from acm2 where id not in (select acmid from acmpart2);
    DECLARE continue handler for not found set c_acm2_finished = 1;

    INSERT INTO migration_status SET description = '20190225_fix_acmpart2_associations repairing acm2
rows that are missing acmpart2 rows';
    ACMPopulate: BEGIN

        -- get list of acm2 rows that dont have acmpart2 rows
        OPEN c_acm2;
        -- with each acm2 row needing fixed (loop)
        get_acm2: LOOP
            -- get row into variables
            FETCH c_acm2 INTO vACMID, vACMName;
            IF c_acm2_finished = 1 THEN
                CLOSE c_acm2;
                LEAVE get_acm2;
            END IF;
            INSERT INTO migration_status SET description = concat('20190225_fix_acmpart2_associations
analyzing acm with id ', vACMID);
            -- with each part (loop)
```



```

get_acmnamepart2: LOOP
-- check if still processing
IF length(vACMName) = 0 THEN
    LEAVE get_acmnamepart2;
END IF;
-- get next part as delineated by semicolon
SET vKeyValuePart := substring_index(substring_index(vACMName, ';', 1), ';', -1);
-- remove it from the head along with trailing semicolon leaving the tail to process on
next loop

SET vACMName := substr(vACMName, length(vKeyValuePart)+2);
-- get key name
SET vKeyName := substring_index(substring_index(vKeyValuePart, '=', 1), '=', -1);
-- remove key from the part with trailing equal sign leaving only the values
SET vKeyValuePart := substr(vKeyValuePart, length(vKeyName)+2);
-- insert the key name and get its id
IF (SELECT 1=1 FROM acmkey2 WHERE name = vKeyName) IS NULL THEN
    INSERT INTO migration_status SET description =
concat('20190225_fix_acmpart2_associations adding missing acm key ', vKeyName, ' to acmkey2');
    INSERT INTO acmkey2 (name) VALUES (vKeyName);
    SET vKeyID := LAST_INSERT_ID();
ELSE
    SELECT id INTO vKeyID FROM acmkey2 WHERE name = vKeyName LIMIT 1;
END IF;
-- with each key/value (loop)
get_acmvalue2: LOOP
-- check if still processing values
IF length(vKeyValuePart) = 0 THEN
    LEAVE get_acmvalue2;
END IF;
-- get next value as delineated by comma
SET vValueName := substring_index(substring_index(vKeyValuePart, ',', 1), ',', -1);
-- remove it from the head along with trailing comma leaving the tail to process on
next loop

SET vKeyValuePart := substr(vKeyValuePart, length(vValueName)+2);
-- insert the value name and get its id
IF (SELECT 1=1 FROM acmvalue2 WHERE name = vValueName) IS NULL THEN
    INSERT INTO migration_status SET description =
concat('20190225_fix_acmpart2_associations adding missing acm value ', vValueName, ' to acmvalue2');
    INSERT INTO acmvalue2 SET name = vValueName;
    SET vValueID := LAST_INSERT_ID();
ELSE
    SELECT id INTO vValueID FROM acmvalue2 WHERE name = vValueName LIMIT 1;
END IF;
-- create association row between acm, key and value
IF (SELECT 1=1 FROM acmpart2 WHERE acmid = vACMID and acmkeyid = vKeyID and
acmvalueid = vValueID) IS NULL THEN
    INSERT INTO migration_status SET description =
concat('20190225_fix_acmpart2_associations for acm with id ', vACMID, ': associating key ',vKeyName,'
with value ',vValueName);
    INSERT INTO acmpart2 (acmid, acmkeyid, acmvalueid) VALUES (vACMID, vKeyID,
vValueID);
    SET vPartID := LAST_INSERT_ID();
ELSE
    SELECT id INTO vPartID FROM acmpart2 WHERE acmid = vACMID and acmkeyid = vKeyID
and acmvalueid = vValueID LIMIT 1;
END IF;
-- end loop processing values
END LOOP get_acmvalue2;
-- end loop for getting part
END LOOP get_acmnamepart2;
-- end loop for acm2
END LOOP get_acm2;
END ACMPOPULATE;
INSERT INTO migration_status SET description = '20190225_fix_acmpart2_associations data repairing
acm2 rows that are missing acmpart2 rows is done.';

```

```
END;  
//  
DELIMITER ;  
CALL sp_Patch_20190225_fix_acmpart2_associations();  
DROP PROCEDURE IF EXISTS sp_Patch_20190225_fix_acmpart2_associations;
```

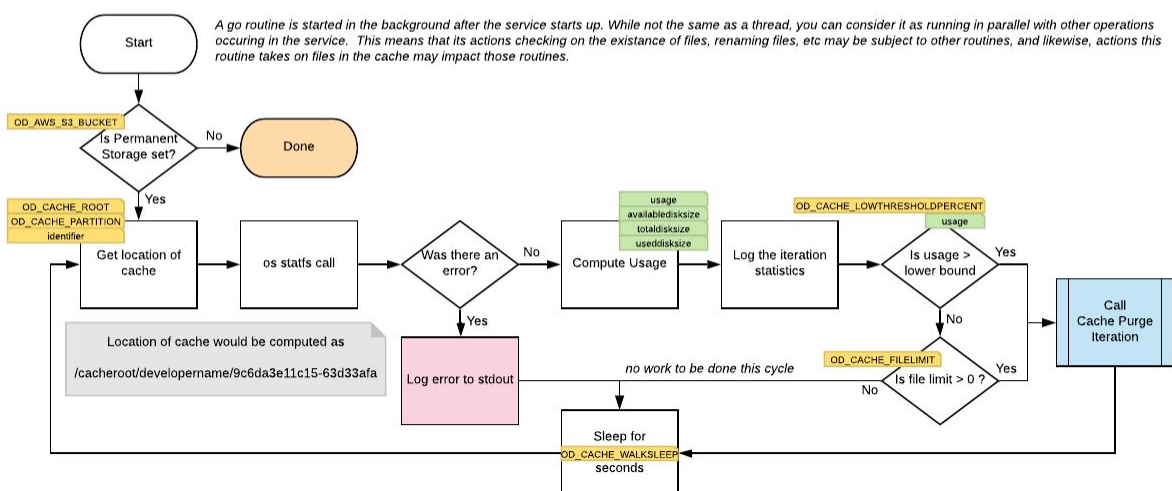
File Cache Management

The local file cache when used in conjunction with permanent storage in S3 has a background thread that monitors the cache and clears space over time as needed to keep within the allocated disk space percentage. Two new features added in 1.0.20 include limiting the number of files, and configuring a time for the process to go to sleep between file checks releasing control via the runtime scheduler to other threads. This section provides a high level overview of the logic flow for how the cache is managed and where the configuration settings factor in

Cache Purge Cycle

The main routine is started when the service itself starts up. If there is no permanent storage configured, it will immediately exit and there will be no cache management as files are not safely evacuated anywhere. For each cycle, a check of the current disk utilization is performed to determine whether any work may potentially need to be done. If the system is configured to allow an unlimited number of files, and the disk usage percentage is below the low threshold, then this cycle is complete. Otherwise, it will go into the cache purge iteration

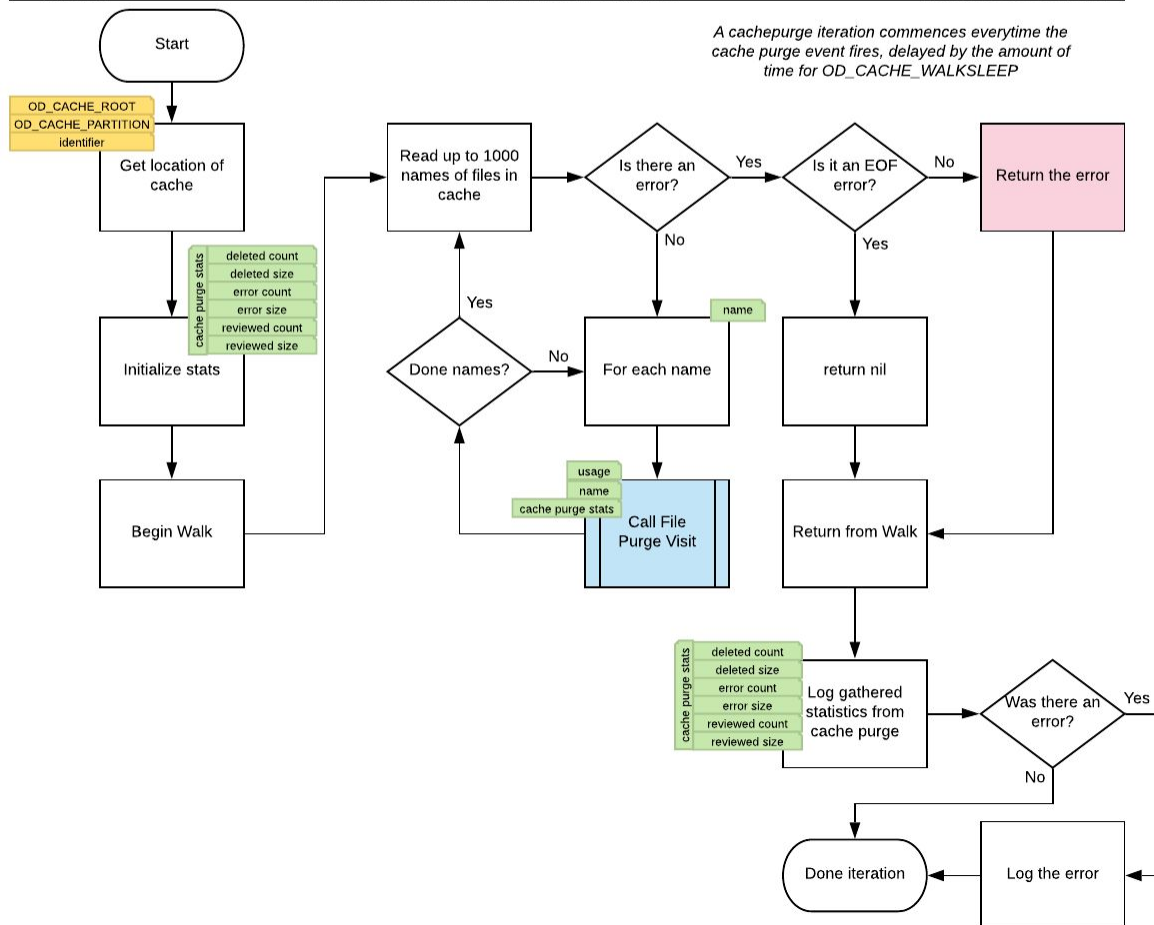
Inputs	Environment Variables	Sample Values
	OD_AWS_S3_BUCKET name of the s3 bucket that acts as permanent file storage OD_CACHE_EVICTAGE number of seconds before file can be deleted from cache OD_CACHE_FILELIMIT quantity of files in cache that can trigger fullness OD_CACHE_FILESLEEP number of milliseconds to wait between each file in cache is processed OD_CACHE_HIGHTHRESHOLDPERCENT high percentage of disk volume where must evict files regardless of age OD_CACHE_LOWTHRESHOLDPERCENT low percentage of disk volume that triggers fullness OD_CACHE_PARTITION optional filepath element to cache, primarily used for multitenancy OD_CACHE_ROOT the root of the cache as an absolute path OD_CACHE_WALKSLEEP number of seconds to wait between cache check iterations	s3cachebucket 120 1000 1 80 40 developername /cacheroor 30
Database		
identifier	assigned on initialization of database for uniqueness	9c6da3e11c15-63d33afa



Cache Purge Iteration

A cache purge iteration will walk each file within the cache, calling the File Purge Visit for the file and logging a debug statement that includes the resulting statistics from the iteration as a whole. If you want to see the total number of files checked, count of errors and how many files were deleted and total size in bytes of those files, then set LOG_LEVEL to DEBUG

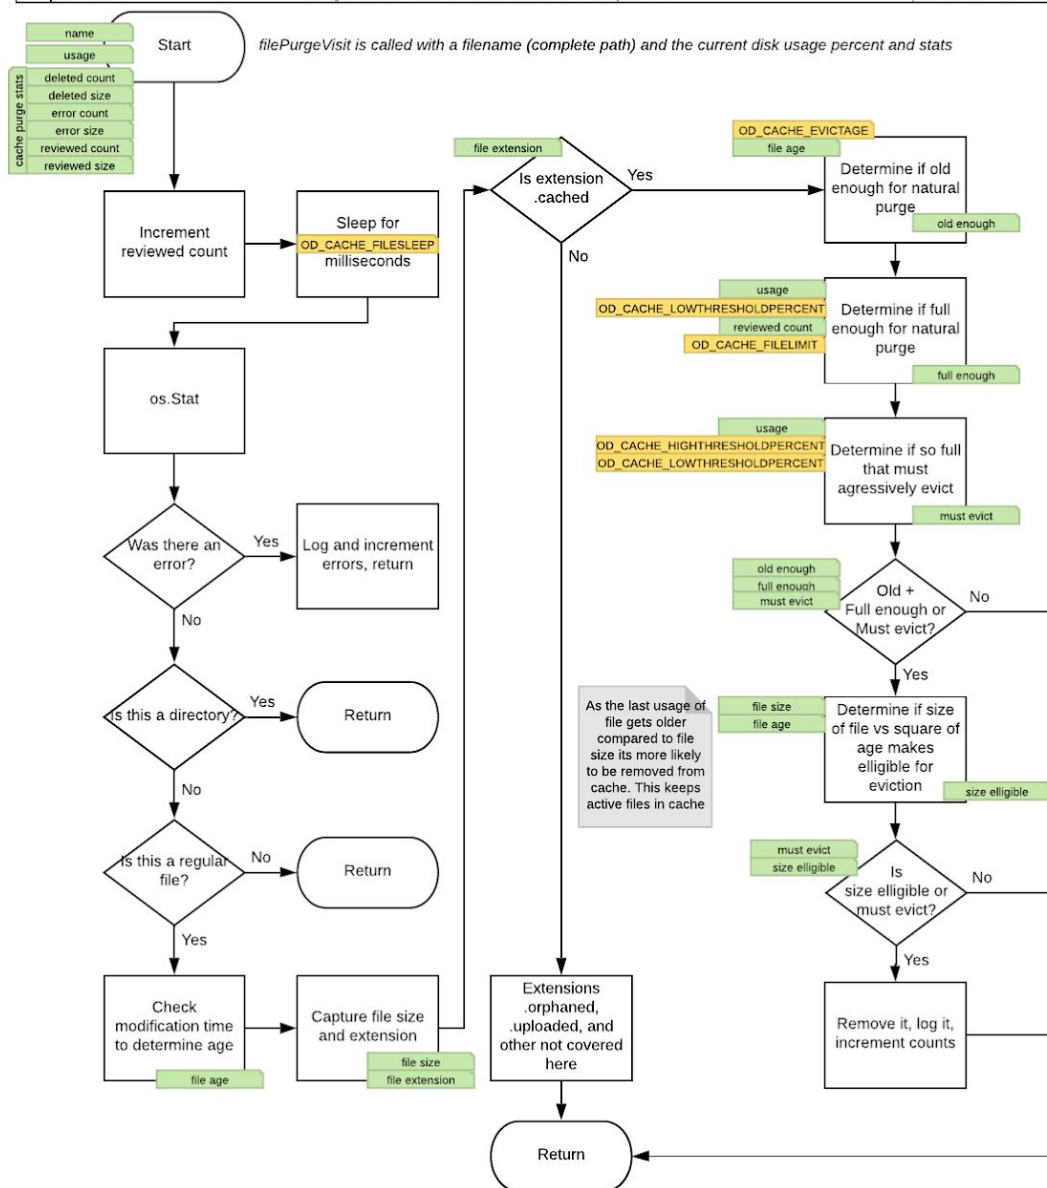
Environment Variables		Sample Values
Inputs	OD_AWS_S3_BUCKET	name of the s3 bucket that acts as permanent file storage
	OD_CACHE_EVICTAGE	number of seconds before file can be deleted from cache
	OD_CACHE_FILELIMIT	quantity of files in cache that can trigger fullness
	OD_CACHE_FILESLEEP	number of milliseconds to wait between each file in cache is processed
	OD_CACHE_HIGHTHRESHOLDPERCENT	high percentage of disk volume where must evict files regardless of age
	OD_CACHE_LOWTHRESHOLDPERCENT	low percentage of disk volume that triggers fullness
	OD_CACHE_PARTITION	optional filepath element to cache, primarily used for multitenancy
	OD_CACHE_ROOT	the root of the cache as an absolute path
Database		
OD_CACHE_WALKSLEEP		number of seconds to wait between cache check iterations
Database		
identifier	assigned on initialization of database for uniqueness	9c6da3e11c15-63d33afa



File Purge Visit

This routine is where all the logic happens on an individual file basis.

Inputs	Environment Variables	Sample Values
	OD_AWS_S3_BUCKET OD_CACHE_EVICTAGE OD_CACHE_FILELIMIT OD_CACHE_FILESLEEP OD_CACHE_HIGHTHRESHOLDPERCENT OD_CACHE_LOWTHRESHOLDPERCENT OD_CACHE_PARTITION OD_CACHE_ROOT OD_CACHE_WALKSLEEP Database	s3cachebucket 120 1000 1 80 40 developername /cacheroot 30 identifier
	assigned on initialization of database for uniqueness	9c6da3e11c15-63d33afa



At the outset, the routine can be configured to sleep for a period of milliseconds for each file. This is helpful for very large caches where otherwise this process may be very CPU and Disk IO intensive resulting in potential delays in the service elsewhere. If you are seeing periods of high CPU utilization lasting for several seconds, consider setting `OD_CACHE_FILESLLEEP` to a value greater than 0. Values of 1 or 2 are commended. A millisecond is a very long time and depending on the quantity of files in cache, can result in much longer periods of processing the iteration. From a resource perspective, using hardware with dedicated IOPS for disk can greatly improve the speed of this routine.

The `os.stat` call performed is also CPU and Disk intensive but is necessary for each file to determine its size, its age of last modification, and whether it is a regular file or directory. Only regular files, with a file extension of `".cached"` are covered here. There are other file cache states but the file purge visit will only remove a file if it has an extension of `".cached"`, indicating that it has been safely evacuated to permanent storage.

The `OD_CACHE_EVICTAGE` defines a preferred minimum age for which files may live in cache before they would be considered for deletion. Typically this is set to a few minutes, but acts as an additional control to balance deletions against the need to re-retrieve file from S3 on a cache miss.


If the current disk usage is over the `OD_CACHE_LOWTHRESHOLDPERCENT` or the number of files reviewed exceeds the `OD_CACHE_FILELIMIT`, then the cache is considered full enough for evicting files. A comparison is also done for usage over the `OD_CACHE_HIGHTHRESHOLDPERCENT`, in which case an eviction must occur to free up space to get back under the threshold.

The age in seconds of a file is compared to the size of that file in such a way as to allow larger files to live in the cache longer before they are deleted. The formula for this is the size of the file, divided by the age of the file in seconds squared, rounded down. If the result is equal to zero, then the file would be deleted.

For example, if a file is 400,000 bytes in size, and it hasn't been accessed for 10 minutes (600 seconds), then the score = $400000 / (600 * 600) \rightarrow 400000 / 360000 \rightarrow 1.111$. The value is greater than 0 so it wouldn't be deleted. However, if it hasn't been accessed in 10 minutes, 33 seconds, then the score = $400000 / (633 * 633) \rightarrow 400000 / 400689 \rightarrow .9982$, rounds down to 0 making it eligible for eviction on the next cache purge.

File Size	Approximate Removal Time After Last Access For Size Eligible Eviction
20 KB	2.5 minutes
400 KB	11 minutes
1 MB	17 minutes
5 MB	37.5 minutes
20 MB	1 hour, 15 minutes
100 MB	2 hours, 47 minutes
1 GB	About 9 hours
4 GB	About 18 hours

The amount of disk space available basically sets an upper bound on the size of a file that can be uploaded to Object Drive at any given time. This is because the file must first be fully received into local cache before being evacuated to permanent storage, and thus eligible for eviction from the cache during the next cache purge cycle.



A bulk migration or upload into Object Drive may necessitate a larger disk size, or more importantly, more free space between cache purge cycles. More aggressive settings may need to be configured for the following

- OD_CACHE_EVICTAGE
- OD_CACHE_LOWTHRESHOLDPERCENT
- OD_CACHE_HIGHTHRESHOLDPERCENT
- OD_CACHE_FILELIMIT
- OD_CACHE_WALKSLEEP

Likewise, a reindexing process can put significant strain on Object Drive with regards to cache misses. If a full reindex is performed for example, every file must be at some point retrieved from permanent storage into the local cache so that its contents can be read. For these unique use cases, it may be prudent to not only leverage more aggressive configuration settings, but use dedicated instances in a cluster to accommodate such bulk migrations or index processes.